

Human-edit steering: 3 of 4 named axes steer the model
in their own vocabulary; the executable rung is certified read-only at $\pm 3\sigma$ and $\pm 6\sigma$ (L5/L6)

edited axis ($+3\sigma$ on the English gloss, re-encoded)

